

Name :- Suraj Singh

CU240251579

BTech CSE 3rd Year SECTION - "B"

Roll No. :- 59

=====

Experiment No. 1

Experiment Name

Real-World Data Collection, Cleaning and Preprocessing

using Python

=====

Aim:

To collect a real-world dataset from publicly available

sources and perform data preprocessing by handling missing

values, duplicate records, inconsistent formats, categorical

variables, and outliers using Python.

Github Link:- [https://github.com/Suraj2007-del/Mastery-in-Data-Analytics-and-Visualizations-and-Career-Advancement-Technical/blob/main/Mastery%20in%20Data%20Analytics%20and%20Visualizations%20and%20Career%20Advancement%20\(Technical\)/Experiment1/Experiment%201.py](https://github.com/Suraj2007-del/Mastery-in-Data-Analytics-and-Visualizations-and-Career-Advancement-Technical/blob/main/Mastery%20in%20Data%20Analytics%20and%20Visualizations%20and%20Career%20Advancement%20(Technical)/Experiment1/Experiment%201.py)



Step 1: Import Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tkinter import Tk, filedialog
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

Step 2: Upload / Select Dataset from Computer

```
import io
from google.colab import files

print("Please select a CSV file to upload.")

uploaded = files.upload()

if not uploaded:
    print("No dataset selected.")
else:
    file_name = next(iter(uploaded)) # Get the name of the first uploaded file
    df = pd.read_csv(io.BytesIO(uploaded[file_name]))

    print("Dataset loaded successfully.")
    print("Selected Dataset:", file_name)
```

Please select a CSV file to upload.
Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving WA_Fn-UseC_-HR-Employee-Attrition.csv to WA_Fn-UseC_-HR-Employee-Attrition.csv
Dataset loaded successfully.
Selected Dataset: WA_Fn-UseC_-HR-Employee-Attrition.csv

Step 3: Display First Five Records

```
print("\n===== FIRST FIVE RECORDS =====")

print(df.head())
```

```
===== FIRST FIVE RECORDS =====
   Age Attrition   BusinessTravel  DailyRate   Department \
0    41      Yes      Travel_Rarely    1102      Sales
1    49      No   Travel_Frequently     279  Research & Development
2    37      Yes      Travel_Rarely    1373  Research & Development
3    33      No   Travel_Frequently    1392  Research & Development
4     27      No      Travel_Rarely     591  Research & Development

   DistanceFromHome  Education  EducationField  EmployeeCount  EmployeeNumber \
0                1         2      Life Sciences              1                1
1                8         1      Life Sciences              1                2
2                2         2         Other              1                4
3                3         4      Life Sciences              1                5
4                2         1         Medical              1                7

   ... RelationshipSatisfaction  StandardHours  StockOptionLevel \
0 ...                1                80                0
1 ...                4                80                1
2 ...                2                80                0
3 ...                3                80                0
4 ...                4                80                1

   TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance  YearsAtCompany \
0                8                0                1                6
1               10                3                3               10
2                7                3                3                0
3                8                3                3                8
4                6                3                3                2

   YearsInCurrentRole  YearsSinceLastPromotion  YearsWithCurrManager
0                4                0                5
1                7                1                7
2                0                0                0
3                7                3                0
4                2                2                2
```

[5 rows x 35 columns]

Step 4: Display Dataset Information

```
print("\n===== DATASET INFORMATION =====")

print(df.info())
```

```
===== DATASET INFORMATION =====
<class 'pandas.core.frame.DataFrame' >
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   Age                 1470 non-null  int64  
 1   Attrition           1470 non-null  object  
 2   BusinessTravel      1470 non-null  object  
 3   DailyRate           1470 non-null  int64  
 4   Department          1470 non-null  object  
 5   DistanceFromHome    1470 non-null  int64  
 6   Education            1470 non-null  int64  
 7   EducationField       1470 non-null  object  
 8   EmployeeCount        1470 non-null  int64  
 9   EmployeeNumber       1470 non-null  int64  
10   EnvironmentSatisfaction 1470 non-null  int64  
11   Gender              1470 non-null  object  
12   HourlyRate          1470 non-null  int64  
13   JobInvolvement       1470 non-null  int64  
14   JobLevel            1470 non-null  int64  
15   JobRole             1470 non-null  object  
16   JobSatisfaction      1470 non-null  int64  
17   MaritalStatus        1470 non-null  object  
18   MonthlyIncome        1470 non-null  int64  
19   MonthlyRate          1470 non-null  int64  
20   NumCompaniesWorked   1470 non-null  int64  
21   Over18              1470 non-null  object  
22   OverTime            1470 non-null  object  
23   PercentSalaryHike    1470 non-null  int64  
24   PerformanceRating    1470 non-null  int64  
25   RelationshipSatisfaction 1470 non-null  int64  
26   StandardHours        1470 non-null  int64  
27   StockOptionLevel     1470 non-null  int64  
28   TotalWorkingYears    1470 non-null  int64  
29   TrainingTimesLastYear 1470 non-null  int64  
30   WorkLifeBalance      1470 non-null  int64  
31   YearsAtCompany       1470 non-null  int64  
32   YearsInCurrentRole   1470 non-null  int64  
33   YearsSinceLastPromotion 1470 non-null  int64  
34   YearsWithCurrManager 1470 non-null  int64  
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
None
```

Step 5: Display Statistical Description

```
print("\n===== STATISTICAL DESCRIPTION =====")

print(df.describe(include="all"))
```

```
===== STATISTICAL DESCRIPTION =====
count    1470.000000    1470    1470    1470.000000 \
unique         NaN         2         3         NaN
top           NaN         No   Travel_Rarely         NaN
freq          NaN        1233        1043         NaN
mean         36.923810         NaN         NaN    802.485714
std           9.135373         NaN         NaN    403.509100
min          18.000000         NaN         NaN    102.000000
25%          30.000000         NaN         NaN    465.000000
50%          36.000000         NaN         NaN    802.000000
75%          43.000000         NaN         NaN   1157.000000
max          60.000000         NaN         NaN   1499.000000

count         Department    DistanceFromHome    Education    EducationField \
unique          3         NaN         NaN         NaN         6
top   Research & Development         NaN         NaN    Life Sciences
freq          961         NaN         NaN         NaN         606
mean          NaN         9.192517         2.912925         NaN
std          NaN         8.106864         1.024165         NaN
min          NaN         1.000000         1.000000         NaN
25%          NaN         2.000000         2.000000         NaN
50%          NaN         7.000000         3.000000         NaN
75%          NaN        14.000000         4.000000         NaN
max          NaN        29.000000         5.000000         NaN

count    EmployeeCount    EmployeeNumber    ...    RelationshipSatisfaction \
unique         NaN        1470.000000    ...         1470.000000
top         NaN         NaN    ...         NaN
freq         NaN         NaN    ...         NaN
mean         1.0        1024.865306    ...         2.712245
std         0.0        602.024335    ...         1.081209
min         1.0         1.000000    ...         1.000000
25%         1.0        491.250000    ...         2.000000
50%         1.0        1020.500000    ...         3.000000
75%         1.0        1555.750000    ...         4.000000
max         1.0        2068.000000    ...         4.000000

count    StandardHours    StockOptionLevel    TotalWorkingYears \
unique         NaN         NaN         NaN
top         NaN         NaN         NaN
freq         NaN         NaN         NaN
mean         80.0         0.793878         11.279592
std         0.0         0.852077         7.788782
min         80.0         0.000000         0.000000
25%         80.0         0.000000         0.000000
50%         80.0         1.000000         10.000000
75%         80.0         1.000000         15.000000
max         80.0         3.000000         40.000000

count    TrainingTimesLastYear    WorkLifeBalance    YearsAtCompany \
unique         NaN         NaN         NaN
top         NaN         NaN         NaN
```

Step 6: Check Dataset Shape

```
print("\n===== DATASET SHAPE =====")

print("Number of Rows:", df.shape[0])
print("Number of Columns:", df.shape[1])
```

```
===== DATASET SHAPE =====
Number of Rows: 1470
Number of Columns: 35
```

Step 7: Identify Missing Values

```
print("\n===== MISSING VALUES =====")

missing_values = df.isnull().sum()

print(missing_values)
```

```
===== MISSING VALUES =====
Age                                0
Attrition                          0
BusinessTravel                     0
DailyRate                          0
Department                         0
DistanceFromHome                   0
Education                          0
EducationField                      0
EmployeeCount                       0
EmployeeNumber                      0
EnvironmentSatisfaction             0
Gender                              0
HourlyRate                          0
JobInvolvement                     0
JobLevel                           0
JobRole                            0
JobSatisfaction                     0
MaritalStatus                       0
MonthlyIncome                      0
MonthlyRate                         0
NumCompaniesWorked                 0
Over18                              0
OverTime                           0
PercentSalaryHike                  0
PerformanceRating                  0
RelationshipSatisfaction            0
StandardHours                      0
StockOptionLevel                   0
TotalWorkingYears                  0
TrainingTimesLastYear              0
WorkLifeBalance                    0
YearsAtCompany                     0
YearsInCurrentRole                  0
YearsSinceLastPromotion             0
YearsWithCurrManager                0
dtype: int64
```

Step 8: Handle Missing Values

```
# Numerical columns -> Median
numerical_columns = df.select_dtypes(
    include=np.number
).columns

for column in numerical_columns:
    if df[column].isnull().sum() > 0:
        df[column] = df[column].fillna(
            df[column].median()
        )

# Categorical columns -> Mode
categorical_columns = df.select_dtypes(
    exclude=np.number
).columns

for column in categorical_columns:
    if df[column].isnull().sum() > 0:
        mode_value = df[column].mode()

        if len(mode_value) > 0:
            df[column] = df[column].fillna(
                mode_value[0]
            )

print("\n===== MISSING VALUE HANDLING =====")

print("Numerical missing values filled using Median.")
print("Categorical missing values filled using Mode.")

print("\nRemaining Missing Values:")

print(df.isnull().sum().sum())
```

```
===== MISSING VALUE HANDLING =====
Numerical missing values filled using Median.
Categorical missing values filled using Mode.

Remaining Missing Values:
0
```

Step 9: Detect Duplicate Records

```
duplicate_count = df.duplicated().sum()

print("\n===== DUPLICATE RECORDS =====")

print("Number of Duplicate Records:", duplicate_count)

# Remove duplicates
df = df.drop_duplicates()

print("Duplicate Records Removed Successfully.")
print("Remaining Rows:", len(df))
```

```
===== DUPLICATE RECORDS =====
Number of Duplicate Records: 0
Duplicate Records Removed Successfully.
Remaining Rows: 1470
```

Step 10: Standardize Categorical Values

```
print("\n===== CATEGORICAL DATA =====")

categorical_columns = df.select_dtypes(
    exclude=np.number
).columns

for column in categorical_columns:

    # Remove unnecessary spaces
    df[column] = df[column].astype(str).str.strip()

    # Convert text values to consistent lowercase format
    df[column] = df[column].str.lower()

print("Categorical values standardized successfully.")
```

```
===== CATEGORICAL DATA =====
Categorical values standardized successfully.
```

Step 11: Encode Categorical Variables

```
encoded_df = df.copy()

categorical_columns = encoded_df.select_dtypes(
    exclude=np.number
).columns

label_encoder = LabelEncoder()

for column in categorical_columns:

    encoded_df[column] = label_encoder.fit_transform(
        encoded_df[column].astype(str)
    )

print("Categorical variables encoded successfully.")
print("Encoding Method: Label Encoding")
```

```
Categorical variables encoded successfully.
Encoding Method: Label Encoding
```

Step 12: Detect Outliers using IQR Method

```
print("\n===== OUTLIER DETECTION =====")

numeric_columns = encoded_df.select_dtypes(
    include=np.number
).columns

outlier_summary = {}

for column in numeric_columns:

    Q1 = encoded_df[column].quantile(0.25)
    Q3 = encoded_df[column].quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = encoded_df[
        (encoded_df[column] < lower_bound) |
        (encoded_df[column] > upper_bound)
    ]

    outlier_summary[column] = len(outliers)

print(
    column,
    "-> Outliers:",
    len(outliers)
)
```

```
===== OUTLIER DETECTION =====
Age -> Outliers: 0
Attrition -> Outliers: 237
BusinessTravel -> Outliers: 0
DailyRate -> Outliers: 0
Department -> Outliers: 0
DistanceFromHome -> Outliers: 0
Education -> Outliers: 0
EducationField -> Outliers: 0
EmployeeCount -> Outliers: 0
EmployeeNumber -> Outliers: 0
EnvironmentSatisfaction -> Outliers: 0
Gender -> Outliers: 0
HourlyRate -> Outliers: 0
JobInvolvement -> Outliers: 0
JobLevel -> Outliers: 0
JobRole -> Outliers: 0
JobSatisfaction -> Outliers: 0
MaritalStatus -> Outliers: 0
MonthlyIncome -> Outliers: 114
MonthlyRate -> Outliers: 0
NumCompaniesWorked -> Outliers: 52
Over18 -> Outliers: 0
OverTime -> Outliers: 0
PercentSalaryHike -> Outliers: 0
PerformanceRating -> Outliers: 226
RelationshipSatisfaction -> Outliers: 0
StandardHours -> Outliers: 0
StockOptionLevel -> Outliers: 85
TotalWorkingYears -> Outliers: 63
TrainingTimesLastYear -> Outliers: 238
WorkLifeBalance -> Outliers: 0
YearsAtCompany -> Outliers: 104
YearsInCurrentRole -> Outliers: 21
YearsSinceLastPromotion -> Outliers: 107
YearsWithCurrManager -> Outliers: 14
```

Step 13: Visualize Outliers using Box Plots

```

boxplot_columns = encoded_df.select_dtypes(
    include=np.number
).columns

if len(boxplot_columns) > 0:

    plt.figure(figsize=(12, 6))


    sns.boxplot(
        data=encoded_df[boxplot_columns]
    )

    plt.title("Box Plot for Outlier Detection")
    plt.xticks(rotation=45)
    plt.tight_layout()

    plt.show()

print("Box Plot generated successfully.")

```



Step 14: Treat Outliers using IQR Capping

```

for column in numeric_columns:

    Q1 = encoded_df[column].quantile(0.25)
    Q3 = encoded_df[column].quantile(0.75)

    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    encoded_df[column] = encoded_df[column].clip(
        lower_bound,
        upper_bound
    )

print("Outliers treated using IQR Capping.")

Outliers treated using IQR Capping.

```

Step 15: Normalize / Standardize Numerical Attributes

```

numeric_columns = encoded_df.select_dtypes(
    include=np.number
).columns

scaler = StandardScaler()

if len(numeric_columns) > 0:

    encoded_df[numeric_columns] = scaler.fit_transform(
        encoded_df[numeric_columns]
    )

print("\n===== FEATURE SCALING =====")

print("Numerical attributes standardized successfully.")
print("Scaling Method: Standardization")
print("Mean approximately equals 0.")
print("Standard deviation approximately equals 1.")

===== FEATURE SCALING =====
Numerical attributes standardized successfully.
Scaling Method: Standardization
Mean approximately equals 0.
Standard deviation approximately equals 1.

```

Step 16: Feature Engineering

```

if "sibsp" in encoded_df.columns and "parch" in encoded_df.columns:

    encoded_df["family_size"] = (
        encoded_df["sibsp"] +
        encoded_df["parch"] +
        1
    )

```

```

    )

    print("\nFeature Created: family_size")

# If Age is available, create AgeGroup
if "age" in df.columns:

    age_values = df["age"]

    encoded_df["age_group"] = pd.cut(
        age_values,
        bins=[0, 12, 18, 35, 60, 100],
        labels=[
            "child",
            "teenager",
            "adult",
            "middle_aged",
            "senior"
        ]
    )

# Encode Age Group
encoded_df["age_group"] = LabelEncoder().fit_transform(
    encoded_df["age_group"].astype(str)
)

print("Feature Created: age_group")

```

Step 17: Display Preprocessed Dataset

```

print("\n===== PREPROCESSED DATASET =====")

print(encoded_df.head())

```

```

===== PREPROCESSED DATASET =====
Age  Attrition  BusinessTravel  DailyRate  Department \
0  0.446350      0.0      0.590048  0.742527  1.401512
1  1.322365      0.0     -0.913194 -1.297775  -0.493817
2  0.008343      0.0      0.590048  1.414363  -0.493817
3 -0.429664      0.0     -0.913194  1.461466  -0.493817
4 -1.086676      0.0      0.590048 -0.524295  -0.493817

DistanceFromHome  Education  EducationField  EmployeeCount  EmployeeNumber \
0      -1.010909   -0.891688      -0.937414           0.0      -1.701283
1      -0.147150  -1.868426      -0.937414           0.0      -1.699621
2      -0.887515  -0.891688      1.316673           0.0      -1.696298
3      -0.764121  1.061787      -0.937414           0.0      -1.694636
4      -0.887515  -1.868426      0.565311           0.0      -1.691313

...  RelationshipSatisfaction  StandardHours  StockOptionLevel \
0  ...      -1.504178           0.0      -0.978160
1  ...      1.191438           0.0      0.300537
2  ...     -0.658973           0.0      -0.978160
3  ...      0.266233           0.0      -0.978160
4  ...      1.191438           0.0      0.300537

TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance  YearsAtCompany \
0      -0.424999      -2.130765      -2.493820      -0.123832
1      -0.151076      0.278639      0.338096      0.678078
2      -0.561960      0.278639      0.338096     -1.326697
3      -0.424999      0.278639      0.338096      0.277123
4      -0.698921      0.278639      0.338096     -0.925742

YearsInCurrentRole  YearsSinceLastPromotion  YearsWithCurrManager
0      -0.058474      -0.776139           0.253823
1      0.785616      -0.372486           0.822605
2     -1.183926      -0.776139          -1.168130
3      0.785616      0.434819          -1.168130
4     -0.621200      0.031166          -0.599348

[5 rows x 35 columns]

```

Step 18: Display Final Dataset Information

```

print("\n===== FINAL DATASET INFORMATION =====")

print("Rows:", encoded_df.shape[0])
print("Columns:", encoded_df.shape[1])

print("\nMissing Values:")
print(encoded_df.isnull().sum().sum())

print("\nDuplicate Records:")
print(encoded_df.duplicated().sum())

```

```

===== FINAL DATASET INFORMATION =====
Rows: 1470
Columns: 35

Missing Values:
0

Duplicate Records:
0

```

Step 19: Save Cleaned Dataset

```

output_file = "cleaned_preprocessed_dataset.csv"

encoded_df.to_csv(
    output_file,
    index=False
)

print("\n===== DATASET STORAGE =====")

print(
    "Cleaned Dataset Saved Successfully:",
    output_file
)

```

)

```
===== DATASET STORAGE =====
Cleaned Dataset Saved Successfully: cleaned_preprocessed_dataset.csv
```

Step 20: Final Output

```
print("\n===== DATA PREPROCESSING =====")

print("Dataset Collection: Completed")
print("Dataset Loading: Completed")
print("Data Inspection: Completed")
print("Missing Value Handling: Completed")
print("Duplicate Removal: Completed")
print("Categorical Standardization: Completed")
print("Categorical Encoding: Completed")
print("Outlier Detection: Completed")
print("Outlier Treatment: Completed")
print("Feature Standardization: Completed")
print("Feature Engineering: Completed")
print("Cleaned Dataset Export: Completed")

print("\nObservations:")

print("Missing numerical values were handled using the median.")
print("Missing categorical values were handled using the mode.")
print("Duplicate records were removed.")
print("Categorical values were standardized and encoded.")
print("Outliers were detected using the IQR method.")
print("Outliers were treated using IQR capping.")
print("Numerical features were standardized using StandardScaler.")
print("New features were created where applicable.")

print("\nExperiment Completed Successfully.")
```

```
===== DATA PREPROCESSING =====
Dataset Collection: Completed
Dataset Loading: Completed
Data Inspection: Completed
Missing Value Handling: Completed
Duplicate Removal: Completed
Categorical Standardization: Completed
Categorical Encoding: Completed
Outlier Detection: Completed
Outlier Treatment: Completed
Feature Standardization: Completed
Feature Engineering: Completed
Cleaned Dataset Export: Completed

Observations:
Missing numerical values were handled using the median.
Missing categorical values were handled using the mode.
Duplicate records were removed.
Categorical values were standardized and encoded.
Outliers were detected using the IQR method.
Outliers were treated using IQR capping.
Numerical features were standardized using StandardScaler.
New features were created where applicable.

Experiment Completed Successfully.
```

EXPERIMENT 1 - ANSWERS TO QUESTIONS

QUESTION 1: Why is data preprocessing considered one of the most important phases in data analytics?

Data preprocessing is important because real-world data is often incomplete, inconsistent, duplicated and noisy.

Before performing visualization, statistical analysis or machine learning, the data must be cleaned and prepared.

Data preprocessing helps to:

Handle missing values. Remove duplicate records. Correct inconsistent data. Detect and treat outliers. Convert categorical variables into numerical form. Bring numerical features to a common scale.

Proper preprocessing improves the quality and reliability of analytical results.

QUESTION 2: Explain different methods of handling missing values with suitable examples.

Missing values can be handled using several methods.

Mean: Missing numerical values can be replaced by the mean of the available values.

Median: Missing numerical values can be replaced by the median. Median is useful when the data contains outliers.

Mode: Missing categorical values can be replaced by the most frequently occurring value.

Row Removal: Rows containing missing values can be removed when the number of missing records is small.

Column Removal: A column can be removed when a very large portion of its values is missing.

Example:

If the Age column contains missing values:

```
df["Age"] = df["Age"].fillna(df["Age"].median())
```

This replaces missing Age values with the median age.

QUESTION 3: Differentiate between Label Encoding and One-Hot Encoding.

Label Encoding:

Label Encoding converts categories into numerical labels.

Example:

Male → 0 Female → 1

It is useful when categories have an inherent order or when a compact numerical representation is required.

One-Hot Encoding:

One-Hot Encoding creates separate binary columns for each category.

Example:

Gender

Male Female

becomes:

Male Female 1 0 0 1

Main Difference:

Label Encoding→ One numerical value per category. One-Hot Encoding→ Separate binary column for each category.

QUESTION 4: What are outliers? How can they affect analytical results?

Outliers are observations that are significantly different from the majority of values in a dataset.

They can occur because of:

Measurement errors. Data entry errors. Unusual events. Natural variation.

Outliers can affect:

Mean and standard deviation. Statistical analysis. Correlation. Machine learning models. Data visualization.

The IQR method can be used to identify outliers.

$IQR = Q3 - Q1$

$Lower\ Bound = Q1 - 1.5 \times IQR$

$Upper\ Bound = Q3 + 1.5 \times IQR$

QUESTION 5: Explain the difference between normalization and standardization.

Normalization:

Normalization generally scales values to a fixed range, commonly 0 to 1.

A common formula is:

$X_{normalized} = (X - X_{min}) / (X_{max} - X_{min})$

Standardization:

Standardization transforms values so that the feature has approximately mean 0 and standard deviation 1.

Formula:

$X_{standardized} = (X - Mean) / Standard\ Deviation$

Main Difference:

Normalization→ Fixed range, commonly 0 to 1. Standardization→ Mean approximately 0 and standard deviation approximately 1.

QUESTION 6: Why should duplicate records be removed before analysis?

Duplicate records represent repeated observations of the same data.

If duplicates are not removed, they can:

Increase the apparent frequency of certain observations. Produce biased statistical results. Affect averages and distributions. Affect machine learning model training. Increase the apparent size of the dataset.

Removing duplicate records helps ensure that each observation is represented appropriately in the dataset.

QUESTION 7: What is feature engineering? Give two practical examples.

Feature engineering is the process of creating new useful attributes from existing data.

It can help represent information in a form that is more useful for analysis or machine learning.

Example 1:

Family Size can be created from:

$SibSp + Parch + 1$

$FamilySize = SibSp + Parch + 1$

Example 2:

Age Group can be created from Age.

For example:

0-12 → Child 13-18 → Teenager 19-35 → Adult 36-60 → Middle Aged 60+ → Senior

Feature engineering can provide additional useful information without collecting new raw data.

QUESTION 8: Which preprocessing techniques would you apply to the IBM HR Employee Attrition dataset and why?

For the IBM HR Employee Attrition dataset, the following preprocessing techniques can be applied:

Missing Value Handling: Identify and appropriately handle missing values.

Duplicate Removal: Remove duplicate employee records.

Categorical Encoding: Convert categorical attributes such as department, job role and marital status into numerical form.

Outlier Detection: Use box plots and the IQR method for numerical variables.

Feature Scaling: Standardize numerical features such as age, income and years of experience where appropriate.

Feature Engineering: Create useful attributes such as age groups or income categories.

These techniques improve data quality and prepare the employee dataset for statistical analysis and machine learning.

QUESTION 9: How does poor-quality data affect machine learning model performance?

Poor-quality data can negatively affect machine learning models.

Problems such as:

Missing values. Duplicate records. Incorrect values. Outliers. Inconsistent categories. Different numerical scales.

can cause models to learn incorrect or misleading patterns.

This may result in:

Lower prediction accuracy. Poor generalization. Biased predictions. Unstable models. Reduced reliability.

Therefore, high-quality preprocessing is important for building reliable machine learning models.

QUESTION 10: Name any three Python libraries commonly used for data preprocessing.

Three commonly used Python libraries are:

Pandas: Used for loading, cleaning, manipulating and analyzing structured datasets.

NumPy: Used for numerical operations and array processing.

Scikit-learn: Provides preprocessing tools such as LabelEncoder, StandardScaler and other machine learning utilities.

Other useful libraries include Matplotlib and Seaborn for visualization and exploratory data analysis.

